

BEYOND REFERENCE MATCHING: SPECIFICATION-BASED EVALUATION OF GENERATED DSP IMPLEMENTATIONS

Xianghui Meng¹ and Jionghao Lin^{1,2,*}

¹The University of Hong Kong

²Carnegie Mellon University

margretmeng1020@gmail.com

*Corresponding author: jionghao@hku.hk

ABSTRACT

A reference filter is one valid solution of a design specification, not a definition of the solution set. On specification-defined filter-design tasks that admit more than one valid realization, we score a candidate by specification-set membership: $S_t = 1$ if it meets a frozen passband, stopband, and stability contract, independently of whether its coefficients match a canonical design. On four FIR/IIR tasks at $f_s = 8$ kHz, nine generated implementations meet the contract while failing coefficient agreement with `firwin` or `butter` (4/4 tasks, 6/12 cells). The same contracts accept 12/12 distinct valid controls and reject 12/12 mechanism mutants. First-principles windowed-sinc and frequency-sampling designs occupy the same set without library designers; a tighter mask still admits two distinct same-order occupants. Coefficient and response-to-reference oracles reject implementations that satisfy the specification. On these tasks, specification-set membership is the correctness criterion we use, not agreement with one realization. Our code is available at: <https://github.com/XianghuiMeng-1020/dsp-semantic-correctness>.

Index Terms— Filter design, specification testing, generated implementations, FIR, IIR

1. INTRODUCTION

DSP implementations are increasingly written by program synthesizers and large language models [1, 2]. The default evaluation reflex is to compare the emitted coefficients, or the emitted frequency response, with a trusted reference implementation [3, 4]. That comparison is exact when the specification is a singleton: one correct integer delay, one declared Fourier convention, one alias period. It is not exact for classical filter design.

A magnitude mask with a free transition band defines a *feasible set* of filters, not a unique coefficient vector [5, 6, 7, 8]. Windowed truncation [9, 10], frequency sampling [11], equiripple Parks–McClellan design [12, 13, 14, 15], and IIR Butterworth or elliptic approximations [16, 17, 18, 19, 20] can

all meet the same passband, stopband, and stability bounds while remaining far apart as impulse responses. A canonical library call—Hamming `firwin`, or a Butterworth section—is one occupant of that set.

This paper treats generated DSP code as an evaluation setting for that structure, not as a model leaderboard. The methodological claim is

reference agreement $\neq$ DSP correctness when $|\mathcal{V}_t| > 1$.
(1)

We freeze a specification S_t , collect generated and first-principles occupants of $\mathcal{V}_t = \{h : S_t(h) = 1\}$, and compare three oracles: coefficient concordance with one valid reference, constrained-band response similarity to that reference, and specification membership. When the task is specification-defined and $|\mathcal{V}_t| > 1$, membership is the correctness criterion used below.

2. SPECIFICATION-SET EVALUATION

Let t be a design task and h a realized filter (FIR taps, or an IIR pair (b, a)). A frozen contract returns $S_t(h) \in \{0, 1\}$. The feasible set is

$$\mathcal{V}_t = \{h : S_t(h) = 1\}. \quad (2)$$

Any valid reference h_r satisfies $h_r \in \mathcal{V}_t$ and nothing stronger. Write $\tilde{h}$ for the coefficient vector of h , or the concatenation (b, a) , truncated to the shorter length when orders differ. Coefficient concordance with h_r is

$$R_{t,r}(h) = \mathbf{1} \left[\frac{\|\tilde{h} - \tilde{h}_r\|_2}{\|\tilde{h}_r\|_2} \leq \tau_R \right], \quad (3)$$

with frozen $\tau_R = 0.05$. $R_{t,r}(h) = 1$ asks whether h is the same realization as h_r . $S_t(h) = 1$ asks whether h is an admissible realization of t . If $\mathcal{V}_t$ contains two realizations more than τ_R apart, then $S_t = 1$ does not imply $R_{t,r} = 1$.

S_t is a magnitude-mask residual on a 4096-point `freqz` grid at $f_s = 8$ kHz. For each constrained band $[f_0, f_1]$ with

bounds $[\ell, u]$,

$$\rho = \max_{f \in [f_0, f_1]} \max(\rho_\ell(f), \rho_u(f)), \quad (4)$$

where $\rho_\ell = (\ell - |H|)_+ / (u - \ell)$ and $\rho_u = (|H| - u)_+ / (u - \ell)$ with a 10^{-6} floor on the denominator, and S_t uses the worst ρ over bands. FIR contracts pass at residual $\leq 10^{-6}$; the IIR contract also requires every pole radius < 0.999 and passes at residual $\leq 10^{-3}$. Neither floor was fit to generated residuals. The four frozen masks are

$$\begin{aligned} \text{LP } |H| &\in [0.95, 1.05] \text{ on } [0, 800], \\ |H| &\leq 0.05 \text{ on } [2000, 4000]; \end{aligned} \quad (5)$$

$$\begin{aligned} \text{BP } |H| &\leq 0.06 \text{ on } [0, 500] \cup [3200, 4000], \\ |H| &\in [0.95, 1.05] \text{ on } [1500, 2200]; \end{aligned} \quad (6)$$

$$\begin{aligned} \text{BS } |H| &\in [0.95, 1.05] \text{ on } [0, 600] \cup [3000, 4000], \\ |H| &\leq 0.06 \text{ on } [1400, 2200]; \end{aligned} \quad (7)$$

$$\begin{aligned} \text{IIR } |H| &\in [0.90, 1.10] \text{ on } [0, 600], \\ |H| &\leq 0.10 \text{ on } [2400, 4000], \end{aligned} \quad (8)$$

with frequencies in hertz. Transition bands are unconstrained, as in the classical approximation problem [5, 14]. Linear phase is not required.

A tone battery T checks consistency of the same mask on sinusoids: real cosines of length 8192 are passed through `lfilter`, and steady-state Hilbert envelopes are compared with the contract bands. T is not an independent gold oracle. Low-pass probes are 50, 400, 800 (pass), 1400 (transition), and 2000, 3000, 3900 Hz (stop); band-pass, band-stop, and IIR probes are the analogous edges and midpoints of (6)–(8). A chirp envelope is not used as a mask residual (valid controls fail that estimator on stopbands).

- **Oracle A** (reference coefficients): $R_{t,r}$.
- **Oracle B** (response similarity): constrained-band $|H|$ RMSE versus h_r , no larger than the maximum same-order control–control RMSE on that task.
- **Oracle C** (specification membership): S_t .

Oracle B is a DSP-native comparison and remains a comparison to *one* element of $\mathcal{V}_t$. Oracle C is membership in $\mathcal{V}_t$.

3. EXPERIMENTAL DESIGN

Generation used Qwen-2.5-72B-Instruct, Llama-3.3-70B-Instruct, and DeepSeek-Chat at temperature 0.7, with four draws per task×model cell. The prompt is the contract statement. No outcome triggered regeneration.

RQ1 — generated implementations. Canonical references are Hamming `firwin` (FIR) and `butter` (IIR) [3]. Before generation, three independently written valid controls per task (`firwin/remez/firls`; `butter/cheby1/ellip`)

all satisfy S_t . Most same-order control pairs exceed τ_R ; the exceptions are band-stop `firls-remez` ($\ell_2 = 0.005$) and IIR `cheby1-ellip` ($\ell_2 = 0.044$). Both exception pairs remain in $\mathcal{V}_t$. Twelve mechanism mutants (band swap, shifted cutoff, truncated length, unstable pole, Nyquist coded as 1 Hz) all fail S_t . The planned design is $4 \times 3 \times 4 = 48$ generations. Same-order control pairs already realize $|\mathcal{V}_t| > 1$: FIR coefficient ℓ_2 ranges are 0.14–0.27 (LP), 0.26–0.38 (BP), and 0.005–0.051 (BS); IIR (b, a) ℓ_2 ranges are 0.044–0.611. Constrained-band $|H|$ RMSE among those pairs stays below 0.001 (FIR) and 0.019 (IIR); full-grid RMSE is 0.074–0.236 (FIR) because the transition is free. Equal-order FIR controls share group delay.

RQ2 — first-principles occupants. Two textbook FIR constructions are formed with `numpy` only. The Type-I Hamming windowed-sinc low-pass of odd length N and cutoff f_c is

$$h[n] = \frac{w[n]}{\sum_k w[k]} 2f_c \text{sinc}(2f_c(n - M)), \quad (9)$$

$M = (N - 1)/2$, $w[n] = 0.54 - 0.46 \cos(2\pi n/(N - 1))$, f_c in cycles/sample [9, 10]. Band-pass and band-stop designs are complementary differences of (9). Frequency sampling assigns a linear-transition desired magnitude $A[k]$ on the N -point DFT grid, applies the Type-I factor $e^{-j\pi(N-1)k/N}$, and inverts [11]. SciPy design symbols are forbidden in that module. Search is over odd N and cutoff/transition placement until the frozen residual is $\leq 10^{-6}$.

RQ3 — tighter mask. One low-pass analysis mask tightens (5) to $|H| \in [0.99, 1.01]$ on $[0, 800]$ Hz and $|H| \leq 0.01$ on $[1400, 4000]$ Hz (fivefold ripple and stop, twofold narrower free transition). `firwin`, (9), and frequency sampling are compared at a common odd length.

RQ4 — oracle disagreement. On the 14 software-eligible generated implementations, Oracles A–C are compared with each other and with the tone consistency check T .

4. RESULTS

Counts are occupancy under a frozen protocol, not prevalence.

RQ1. Of 48 requests, 20 executed and 14 passed the type/shape screen. Nine of those 14 satisfy S_t ; all nine fail Oracle A (Table 1). The event occupies 4/4 tasks and 6/12 generation cells; execution outcomes are heterogeneous across cells. Distinct occupants include a 128-tap `firwin2` low-pass and a 20-tap `remez` low-pass; two 101-tap band-stops with relative $\ell_2 > 0.05$ versus `firwin`; and elliptic (3,3) and Butterworth (4,4) IIR designs, both stable and in-band. Constrained tones through the nine filters are in-mask on 63/63 probes; the twelve valid controls are in-mask on 90/90. On the specification bands, FIR occupants remain close in $|H|$ to the canonical design (RMSE 0.0007–0.0045). Different-order FIR occupants differ in passband group delay by 13.5–30.5 samples. The elliptic pair is not equivalent

to the Butterworth reference on the mask: constrained-band $|H|$ RMSE versus Butterworth is 0.065 against a control-pair maximum of 0.019. Two of the nine programs have the same length as `firwin`; one of those two still exceeds τ_R (0.057). A post-hoc swap of h_r among the three frozen controls leaves 7/9 discordant with all three; the two same-length band-stops become concordant with `remez/firls` while remaining discordant with `firwin`. The swap relocates a program to a different element of $\mathcal{V}_t$, not out of $\mathcal{V}_t$.

RQ2. Both first-principles methods enter $\mathcal{V}_t$ on all three FIR tasks (6/6 shortest designs; Table 2). Locked to the canonical length (81 low-pass, 101 band-pass and band-stop), all six still satisfy S_t ; five of six exceed τ_R versus `firwin`. Hamming windowed-sinc at 101 taps on the band-stop is the same method as `firwin` (coefficient $\ell_2 = 0.0013$, $\Delta\text{GD} = 0$). Frequency sampling at that length remains in $\mathcal{V}_t$ with coefficient ℓ_2 of 0.18–0.39. Every constrained tone is in-mask. Multiplicity is a property of the mask, not of library-API diversity.

RQ3. On the tighter low-pass, the shortest passing lengths are 43 (`firwin` and windowed-sinc; numerically identical) and 57 (frequency sampling). At common length $N = 57$, both distinct methods satisfy the tight contract, pass 7/7 constrained tones, share a 28-sample group delay, and differ by coefficient $\ell_2 = 0.115$ and full-grid $|H|$ RMSE 0.057, with spec-band RMSE 0.0014 (Table 2). A fivefold / twofold tightening of the frozen numbers does not collapse $\mathcal{V}_t$ to a singleton.

RQ4. Table 3 reports agreement of each oracle with the tone consistency check T (same bands as S_t ; not an external gold). Oracle C and T agree on all 14 eligible implementations (9/9 with $S_t = 1$ are fully in-mask; 5/5 with $S_t = 0$ miss at least one constrained tone), as expected for two readings of one mask. Oracle A agrees with T on 5/14 and rejects every implementation that is in-mask on T (9/9). Oracle B agrees with T on 8/14 and still rejects 6/9 in-mask implementations, namely the different-order FIR designs and the elliptic pair. The three filters that are in-mask on T and pass Oracle B—two same-order band-stops and the Butterworth (4, 4)—still fail Oracle A. A reference protocol therefore accepts 0 of 48 planned generations and 0/4 tasks; the specification protocol accepts 9 generations and 4/4 tasks. On band-pass and IIR cells, ranking by coefficient distance elevates a specification-failing implementation above a specification-valid one. The band-pass case is a 101-tap `firwin2` with coefficient $\ell_2 = 0.401$ and residual 0.352 (tones 5/9) ranked above a 128-tap specification-valid `firwin2` with $\ell_2 = 1.369$ and residual 0. The IIR case is a specification-failing design with $\ell_2 = 0.499$ versus `butter` ranked above elliptic (3, 3) and Butterworth (4, 4) occupants with residual 0. Kendall τ between coefficient distance and specification residual on the 14 eligible rows is 0.38. Oracle B calibration maxima (control-pair constrained $|H|$ RMSE) are 8.6×10^{-4} (LP), 1.0×10^{-3} (BP), 7.4×10^{-4} (BS), and 1.9×10^{-2} (IIR).

Table 1. RQ1 occupancy on the frozen specification-set arm. $S_t = 1$ is contract membership; $R_{t,r} = 0$ is coefficient discordance with the canonical reference.

Quantity	Value
Planned / executed / eligible	48 / 20 / 14
$S_t = 1$	9
$S_t = 1$ and $R_{t,r} = 0$	9
Tasks / sources with a witness / cells	4 / 4 / 2 / 3 / 6 / 12
Pre-generation controls ($S_t = 1$)	12 / 12
Pre-generation mutants ($S_t = 0$)	12 / 12
Constrained tones, $S_t = 1$ / controls	63 / 63 / 90 / 90

Table 2. RQ2 and RQ3. WS: Hamming windowed-sinc. FS: frequency sampling. Same-order lengths are the frozen `firwin` lengths (RQ2) or $N = 57$ (RQ3).

Setting	In $\mathcal{V}_t$	Same-order S_t	$\ell_2 > \tau_R$
RQ2 shortest WS/FS (3 FIR tasks)	6/6	—	6/6
RQ2 same-order WS/FS	—	6/6	5/6
RQ3 tight LP, $N = 57$, WS vs. FS	2/2	2/2	1/1 pair

5. CONCLUSION

When a specification-defined DSP task admits multiple valid realizations, a reference implementation is one point in $\mathcal{V}_t$. Generated filters, first-principles windowed-sinc and frequency-sampling designs, and a tightened magnitude mask all exhibit the same structure: $S_t = 1$ does not imply coefficient or response agreement with a single canonical occupant. Coefficient matching and response-to-reference matching remain useful as realization diagnostics. On such tasks they are not, by themselves, correctness tests. The correctness criterion used here is specification-set membership.

Compliance with ethical standards. No funding was received for conducting this study. The authors have no relevant financial or nonfinancial interests to disclose. The study did not involve human participants, animal subjects, or personally identifiable data.

Table 3. RQ4: agreement with the tone consistency check T ($n = 14$ eligible implementations). T uses the same bands as S_t . $T+\wedge\text{orc}-$ is oracle fail and T pass; $T-\wedge\text{orc}+$ is oracle pass and T fail.

Oracle	Agree with T	$T+\wedge\text{orc}-$	$T-\wedge\text{orc}+$
A (coefficients vs. h_r)	5/14	1.00	0
B (band $ H $ vs. h_r)	8/14	0.67	0
C (specification membership)	14/14	0	0

6. REFERENCES

- [1] M. Chen, J. Tworek, H. Jun, et al., “Evaluating large language models trained on code,” *arXiv preprint arXiv:2107.03374*, 2021.
- [2] J. Liu, C. S. Xia, Y. Wang, and L. Zhang, “Is your code generated by ChatGPT really correct? rigorous evaluation of large language models for code generation,” in *Adv. Neural Information Processing Systems (NeurIPS)*, 2023.
- [3] P. Virtanen, R. Gommers, T. E. Oliphant, et al., “SciPy 1.0: Fundamental algorithms for scientific computing in Python,” *Nature Methods*, vol. 17, pp. 261–272, 2020.
- [4] T. Huuhtanen, J. Itkonen, and C. Lassenius, “Combining algebraic and domain testing to design adequate test cases for signal processing algorithms,” in *Proc. IEEE 8th Int. Conf. Software Testing, Verification and Validation (ICST)*, 2015, pp. 1–10.
- [5] A. V. Oppenheim and R. W. Schaffer, *Discrete-Time Signal Processing*, Prentice Hall, Upper Saddle River, NJ, 3 edition, 2010.
- [6] J. G. Proakis and D. G. Manolakis, *Digital Signal Processing: Principles, Algorithms, and Applications*, Pearson Prentice Hall, Upper Saddle River, NJ, 4 edition, 2007.
- [7] S. K. Mitra, *Digital Signal Processing: A Computer-Based Approach*, McGraw-Hill, New York, NY, 4 edition, 2011.
- [8] L. R. Rabiner and B. Gold, *Theory and Application of Digital Signal Processing*, Prentice Hall, Englewood Cliffs, NJ, 1975.
- [9] J. F. Kaiser, “Nonrecursive digital filter design using the I_0 -sinh window function,” in *Proc. IEEE Int. Symp. Circuits and Systems*, 1974, pp. 20–23.
- [10] F. J. Harris, “On the use of windows for harmonic analysis with the discrete Fourier transform,” *Proc. IEEE*, vol. 66, no. 1, pp. 51–83, Jan. 1978.
- [11] L. R. Rabiner, B. Gold, and C. A. McGonegal, “An approach to the approximation problem for nonrecursive digital filters,” *IEEE Trans. Audio and Electroacoustics*, vol. AU-18, no. 2, pp. 83–106, June 1970.
- [12] T. W. Parks and J. H. McClellan, “Chebyshev approximation for nonrecursive digital filters with linear phase,” *IEEE Trans. Circuit Theory*, vol. CT-19, no. 2, pp. 189–194, Mar. 1972.
- [13] J. H. McClellan, T. W. Parks, and L. R. Rabiner, “A computer program for designing optimum FIR linear phase digital filters,” *IEEE Trans. Audio and Electroacoustics*, vol. AU-21, no. 6, pp. 506–526, Dec. 1973.
- [14] L. R. Rabiner, J. H. McClellan, and T. W. Parks, “FIR digital filter design techniques using weighted Chebyshev approximation,” *Proc. IEEE*, vol. 63, no. 4, pp. 595–610, Apr. 1975.
- [15] O. Herrmann, “On the approximation problem in non-recursive digital filter design,” *IEEE Trans. Circuit Theory*, vol. CT-18, no. 3, pp. 411–413, May 1971.
- [16] S. Butterworth, “On the theory of filter amplifiers,” *Experimental Wireless and the Wireless Engineer*, vol. 7, pp. 536–541, 1930.
- [17] A. G. Constantinides, “Spectral transformations for digital filters,” *Proc. IEE*, vol. 117, no. 8, pp. 1585–1590, Aug. 1970.
- [18] L. B. Jackson, *Digital Filters and Signal Processing*, Springer, Boston, MA, 3 edition, 1996.
- [19] A. Antoniou, *Digital Signal Processing: Signals, Systems, and Filters*, McGraw-Hill, New York, NY, 2006.
- [20] P. P. Vaidyanathan, *Multirate Systems and Filter Banks*, Prentice Hall, Englewood Cliffs, NJ, 1993.